

有纸质笔记

Anisotropic Mesh Adaptation for the Manycore Era

Georgios Rokos


Imperial College London

April 23, 2015

- ▶ Irregular applications:
 - Use of unstructured or completely irregular data
 - Mostly represented as graphs
- ▶ Challenges:
 - Unpredictable memory access patterns
 - Poor data locality
 - Kernels end up being memory-bound rather than compute-bound
 - Data-driven algorithms, hard to extract parallelism
 - Fine-grained parallelism leads to frequent thread synchronisation
- ▶ Mutable dependencies:
 - E.g. *morph algorithms* (Pingali *et al.*)
 - Graph topology is mutated in non-trivial ways
 - Any preprocessing is constantly invalidated


- ▶ Need to study a real-world problem in order to develop techniques for parallelising irregular kernels
- ▶ Unstructured meshes and finite element/volume modelling:
 - Spatial domain discretised into triangles (in this talk we only focus on 2D)
 - Ideal for representing complex geometries (e.g. coastal modelling)
 - Numerical solutions of partial differential equations (PDEs)
- ▶ Mesh adaptivity methods:
 - Allow dynamic control of solution error
 - Keep the resolution in the goldilocks zone - not too high and not too low
 - Minimise computational cost for a specific model accuracy

Example



Example: Detail along the wave front

- Elements are stretched along the direction of the front



- ▶ Initial mesh generated *a priori*:
 - Difficult to generate a mesh that is both efficient and resolves the solution where required
 - Particularly difficult for multi-scale problems
- ▶ Local error estimates
 - Error estimate transformed to a metric tensor field (MTF)
 - Discretised vertex-wise
 - Tensor at some vertex specifies local size and shape of an element containing that vertex which is required to achieve a specific error tolerance
- ▶ Support for anisotropic problems
 - PDE exhibits directional dependencies (desired element size and shape) encoded in a MTF
 - E.g. higher resolution is required perpendicular to a shock front (where flow is more complex) than along the shock

Metric tensor

- ▶ A metric tensor is a symmetric matrix, 2x2 in 2D, 3x3 in 3D
- ▶ Defines length of vectors
- ▶ Allows us to calculate inner products in generalised spaces, in the same way the dot product defines distance in Euclidean space
- ▶ Example in 2D with vertices $V_1(x_1, y_1)$, $V_2(x_2, y_2)$ and edge $\mathbf{E} = (x_0, y_0) = (x_2 - x_1, y_2 - y_1)$
 - Length in Euclidean space given by the dot product:

$$L_{Euclidean} = \| \mathbf{E} \| = \sqrt{\mathbf{E} \cdot \mathbf{E}} = \sqrt{x_0^2 + y_0^2}$$

- Edge length with respect to a metric tensor $\mathbf{M} = \begin{bmatrix} A & B \\ B & C \end{bmatrix}$:


$$\begin{aligned} L_M = \| \mathbf{E} \|_M &= \sqrt{\mathbf{E}^T \mathbf{M} \mathbf{E}} = \sqrt{[x_0 \ y_0] \begin{bmatrix} A & B \\ B & C \end{bmatrix} \begin{bmatrix} x_0 \\ y_0 \end{bmatrix}} = \\ &= \sqrt{x_0^2 A + 2x_0 y_0 B + y_0^2 C} \end{aligned} \quad (1)$$

Element size and shape

- ▶ Metric tensor in the middle of a triangle
 - Linear interpolation of metric tensors at the three vertices
- ▶ Eigenvalue decomposition of a 2D metric tensor:

$$\mathbf{M} = \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^T = \begin{bmatrix} Q_{00} & Q_{01} \\ Q_{10} & Q_{11} \end{bmatrix} \begin{bmatrix} \lambda_0 \\ \lambda_1 \end{bmatrix} \begin{bmatrix} Q_{00} & Q_{10} \\ Q_{01} & Q_{11} \end{bmatrix}$$

- ▶ Each eigenvalue λ_i encodes the required element size in the direction of the corresponding eigenvector Q_i



Adaptive algorithms

► 4 adaptive algorithms



- Coarsening
 - Refinement
 - Swapping
- } h-adaptivity \longrightarrow mesh topology is modified
- Smoothing
- } r-adaptivity \longrightarrow mesh topology is not modified

► Mesh adaptation

- Element quality functional measures 'distance' from ideal element as defined by metric field
- Iterative application of local mesh operations until the quality is within some threshold
- h-adaptivity: morph algorithms


Coarsening

- ▶ Done via edge collapse: vertex V_B collapses onto V_A , removing the dashed edge and the adjacent elements from the mesh (Li et al. 2005)
- ▶ Every vertex is examined to determine onto which neighbour (if any) it can collapse
- ▶ If a vertex is removed the local neighbourhood is modified, so all neighbours are marked for re-examination
⇒ Propagation of coarsening




Refinement

- ▶ Edge and element refinement: long edges are split, leading to 1:2 (bisection), 1:3 or 1:4 (regular refinement) division of elements, which increases local mesh resolution (Li et al. 2005)
- ▶ At first, all edges are visited and long edges are split
- ▶ Next up, elements with a split edge are split according to the number of split edges
- ▶ No need for propagation, just execute refinement kernel again





Swapping

- ▶ Edge swapping: edges shared between two elements can be flipped if the minimum quality of the element pair is raised (Li et al. 2005)
- ▶ Improves mesh quality without increasing the number of elements
- ▶ Once an edge has been flipped, all adjacent edges are marked for re-examination
⇒ Propagation of swapping



Smoothing

- Implemented as optimisation-based vertex smoothing: a vertex u_i is relocated to a new position so that the quality of the worst element among $\{e_{i,0}..e_{i,5}\}$ is maximised (Freitag et al. 1995)
- Linear search problem in the direction of the steepest ascent of the derivative of the quality functional
- Smoothing is propagated




Topological hazards

Example:

- ▶ One thread coarsens edge $V_B V_C$, V_B collapses onto V_C
- ▶ Another thread coarsens edge $V_C V_D$, V_C collapses onto V_D


$\Rightarrow V_B$ collapses onto a vertex (V_C) which is being deleted!



Topological hazards: Mesh colouring

Solution: Mesh colouring

- ▶ Nodes are processed in batches of independent sets
 - Guarantees that adjacent nodes cannot collapse at the same time
- ▶ colouring is in the loop
 - Need it to be fast and use as few colours as possible
- ▶ colouring algorithm by Çatalyürek et al.
 - Based on optimistic/speculative execution
 - We developed an improved version (more on that offline)




Race conditions

Example: updating adjacency lists

- ▶ One thread coarsens edge $V_B V_C$, V_B collapses onto V_C
 - adjacency lists of V_C are modified
 - e.g. V_A must be added to the node-node list of V_C
- ▶ Another thread coarsens edge $V_D V_C$, V_D collapses onto V_C
 - adjacency lists of V_C are modified
 - e.g. V_E must be added to the node-node list of V_C

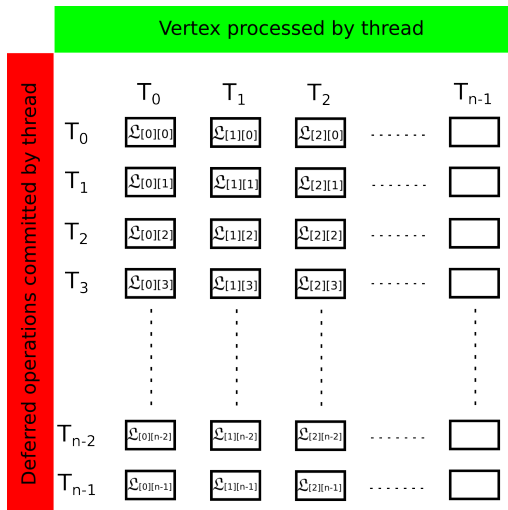
⇒ Both threads try to modify the node-node list of V_C



Race conditions: Deferred updates

Solution: Defer updates until the independent set has been processed

- ▶ Allocate lists $\mathcal{L}_{[i][j]}$ of deferred updates, $i, j = 0..nthreads - 1$
- ▶ A thread T_i stores updates pertaining to vertex V_A in $\mathcal{L}_{[T_i][j]}$, $j = hash(V_A) \% nthreads$
- ▶ At the end, every thread T_j commits all updates in $\mathcal{L}_{[i][T_j]}$, $i=0..N-1$
- ▶ Advantage: Every thread visits only those updates it is responsible for committing \implies FAST!



Worklist: A set of workitems which will be processed, e.g. a global worklist of nodes in an independent set

- ▶ Threads colour the mesh in parallel
 - Every thread stores the nodes it has coloured in local (private) arrays, $local_{[T_i]}[colour]$, $colour=0..ncolours$
 - For each colour C , we need to concatenate all private arrays $local_{[T_i]}[C]$, $i=0..N-1$ into a global array $global_{[C]}$
- ▶ Classic approach: Prefix sum (or “scan” in MPI terminology) on the index in $global_{[C]}$ for every thread
 - Threads need to synchronise \implies SLOW!
- ▶ Alternative: Atomic fetch-and-add
 - Introduced in OpenMP 3.1
 - “atomic capture” directive
 - Older compilers support it either via intrinsics or inline assembly

Worklists: Example

```
1 // Pre-allocate enough space
2 std::vector<Item> globalWorklist(some_appropriate_size);
3 int worklistSize = 0;
4
5 #pragma omp parallel
6 {
7     // Initialise a private list
8     std::vector<Item> private_list;
9
10 #pragma omp for nowait
11     for(all items which need to be processed){
12         do_some_work();
13         private_list.push_back(item);
14     }
15
16     // Private variable – the index in global worklist
17     int idx;
18
19 #pragma omp atomic capture
20     {
21         idx = worklistSize;
22         worklistSize += private_list.size();
23     }
24
25     memcpy(&globalWorklist[idx], &private_list[0], private_list.size() * sizeof(Item));
26 }
```

► Note the “nowait” clause at omp-for

- Threads need not synchronise at the end of the loop \implies FAST!

Loop scheduling: OMP

- ▶ Highly diverse loops.
- ▶ Example: Mesh refinement
 - Element-refinement loop traverses all elements
 - An element can be processed in 4 different ways:
no split, 1:2, 1:3, 1:4

⇒ Load imbalance!
- ▶ OMP dynamic scheduling
 - Perfect load balance
 - Way too much overhead (millions of nodes/elements)

⇒ Poor performance
- ▶ OMP guided scheduling
 - Decent load balance, but it could be better
 - Almost no overhead

⇒ Much better performance

Loop scheduling: Work-stealing

- ▶ Work-stealing scheduler
 - Very good load balance
 - Relatively little overhead

⇒ Work-stealing is the way to go!
- ▶ OMP does not support work-stealing:
 - We had to implement it manually
- ▶ Hand-written scheduler implements an improved version of the classic work-stealing algorithm:
 - Excellent load balance
 - Very little overhead

⇒ Best performance
- ▶ Work on this scheduler is still in progress:
 - Preliminary results from synthetic benchmarks: outperforms Intel[®] Cilk[™] Plus work-stealing

► Parallel anisotropic Adaptive Mesh Toolkit:

- 2D/3D mesh adaptivity framework
- Open source, under the BSD license
- Available on Github

<https://github.com/meshadaptation/pragmatic>


- Implements all aforementioned adaptive algorithms
- Hybrid OpenMP/MPI support
- Currently being integrated with Dolfin (FEniCS) and DMPlex (PETSc)

PRAgMaTlc: Sample benchmark

A synthetic solution ψ is defined to vary in time and space for some value of the period T :


$$\psi(x, y, t) = 0.1 \sin \left(50x + \frac{2\pi t}{T} \right) + \arctan \left(-\frac{0.1}{2x - \sin \left(5y + \frac{2\pi t}{T} \right)} \right)$$

Benchmark solution field for some time step t_i




Sample benchmark: Initial mesh

Initial, auto-generated mesh




Sample benchmark: Adapted mesh snapshot

Adapted mesh for time step t_i




Sample benchmark: Mesh quality snapshot

Quality of adapted mesh for time step t_i




Sample benchmark: Mesh quality detail

Detail of quality around the sinusoidal front



Sample benchmark: Aggregated mesh quality

Aggregated histogram of element quality over all time steps




- ▶ Average element quality: > 0.9 (close to ideal 1.0)
- ▶ Worst element quality: > 0.6

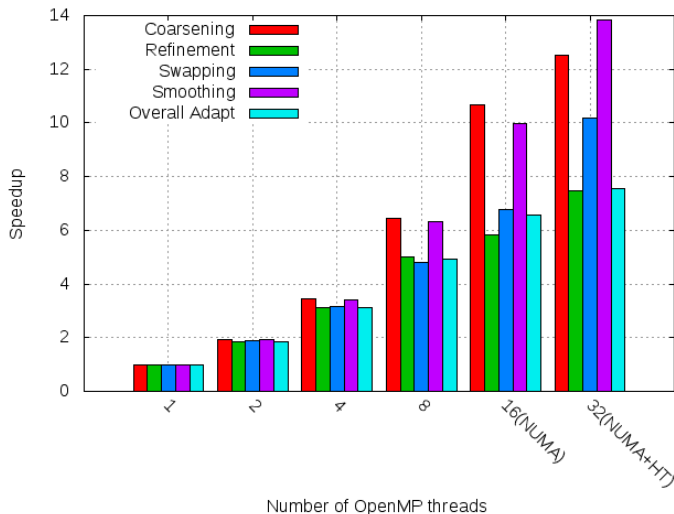
Same sample benchmark

- ▶ $\times 100$ finer metric tensor field, $\approx 500k$ elements, $\approx 250k$ nodes
- ▶ Compiled with Intel[®] Compiler Suite 14.0.1, `-Ofast` flag
- ▶ Executed on a dual-socket Xeon[®] E5-2650 system (Sandy Bridge, 2GHz, 8 cores/16 HT per socket), using thread-core affinity support
- ▶ Execution time over all time steps for:
 - (1) each of the four adaptive algorithms
 - (2) total adapt = sum of the four adaptive algorithms + mesh defragmentation
- ▶ $\approx 1.5s$ per time step with 32 threads
- ▶ low compared with typical solution times

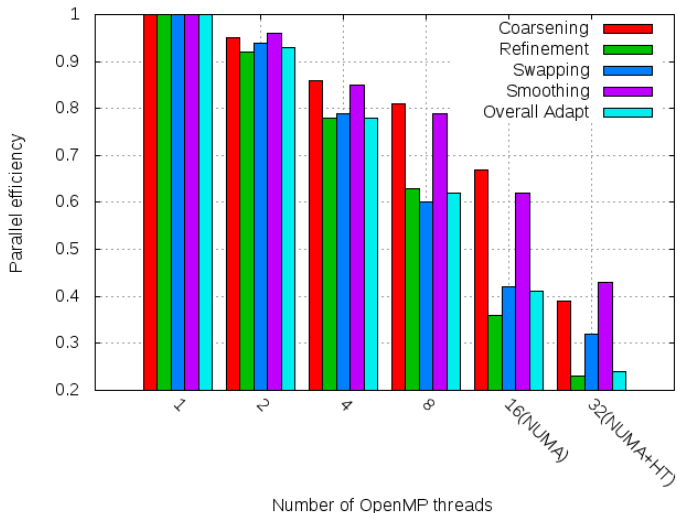
Performance results: execution time



Performance results: speedup



Performance results: parallel efficiency



- ▶ Coarsening and smoothing scale well
 - Scalability is mostly limited by thread synchronisation at the end of every independent set
- ▶ Refinement and swapping are further affected by bandwidth saturation
 - Enabling hyperthreading improves performance considerably
 - Bandwidth saturation is only to be expected for an application with little data locality

Can we do better?

- ▶ Thread synchronisation is the main factor limiting parallel scalability
- ▶ colouring and the deferred-operations mechanism involve thread synchronisation
- ▶ Alternative: optimistic execution
 - Inspired by the Galois framework (Pingali *et al.*)
 - Lock associated with every mesh vertex
 - A thread tries to acquire the locks of all vertices in a local mesh patch
 - If one of the locks is already held by another thread, abort
 - Early experimentation: abort ratio $< 0.01\%$
 - Single-threaded execution is slower (acquiring/releasing locks is expensive)
 - But code becomes more scalable (Pingali reports parallel efficiency of $> 70\%$ on a 512-core SGI Ultraviolet system)

- ▶ PRAgMaTlc produces high-quality adapted meshes
- ▶ Anisotropic mesh adaptivity sounds expensive and hard to parallelise
- ▶ It can be fast enough to pay off in common usage scenarios
- ▶ Some remaining thread synchronisation and bandwidth saturation are currently the limiting factors
- ▶ Current focus is on performance optimisation for 3D and MPI
- ▶ Inherent difficulty of parallelising complex, irregular algorithms:
 - Optimistic colouring, deferred operations, worklists, work-stealing scheduler proved to be keys to high performance
 - This irregular compute methodology can be used in other applications with mutable irregular data

Acknowledgements and further reading

PRAgMaTlc is brought to you by (alphabetically):

- ▶ **Dr. Gerard J. Gorman**, g.gorman@imperial.ac.uk, Department of Earth Science and Engineering, Imperial College London, UK
- ▶ **Prof. Paul H. J. Kelly**, p.kelly@imperial.ac.uk, Department of Computing, Imperial College London, UK
- ▶ **Georgios Rokos**, georgios.rokos09@imperial.ac.uk, Department of Computing, Imperial College London, UK

This project has been supported by Fujitsu Laboratories of Europe Ltd and the EPSRC (grant numbers EP/I00677X/1 and EP/L000407/1).